Gil Raphaelli | g@raphaelli.com | 610-256-4055 | https://github.com/graphaelli

Summary: Developer with strong systems background experienced in building and scaling data intensive applications. Experienced in most phases of product development - writing software with Go and Python to solve hard problems, operating large deployments, managing people and projects, developing product roadmaps, and more.

RealScout | Data Engineering Manager | September 2015 - Present
- Implemented AWS Kinesis/Lambda/ECR based MLS import system to replace existing infrastructure, reducing time to integrate MLSes from months to days while improving accuracy, resiliency, and debuggability.
- Standardized development platform using docker containers.
- Parallelized Caffe-based property photo classification system.
- Implemented various customer facing product features using Ruby/Rails.
- Led team of 4 engineers and data scientists, providing technical guidance and personal mentorship.

Monetate | Software Engineer, Infrastructure Team | February 2015 - September 2015
- Founding member of team to dedicated to application infrastructure, tools, monitoring, and performance
- Built backend for customer facing ad-hoc analytics platform using Druid.
- Replaced home-grown configuration management and deployment system with one utilizing packer, ansible and a lightweight Django app

Monetate | Engineering Director | August 2013-February 2015
- Managed teams of 6-9 developers within product development team
- Delivered on various product initiatives - initially small, eventually most central/complex to organization
- Continued in oncall rotation until backfill hired and trained

Monetate | Systems Engineer | May 2012-August 2013
- Implemented distributed fact collection and processing system, replacing Hive with Amazon Redshift
- Automated build, test, deployment processes increasing confidence in releases.
  - Automated time intensive and error prone application build process.
  - Enhanced homegrown configuration management tools, extensively utilizing AWS apis with boto.
  - Implemented tools for performing differential analysis between deployments.
- Built tools for and performed full-scale testing of seasonal peak load, implemented fixes based on results.

Flickr | Senior Backend Developer | 2010-2012
- Overhauled photo storage backend in Java.
  - Updated networking core while preserving client protocol.
  - Eliminated application performance issues related to underlying storage issues.
  - Simplified code, wrote documentation, exposed more/better metrics versus old system.
  - Wrote end-to-end integration tests in Python.
- Redesigned and wrote search feeding system in Java to improve resource utilization, speed, and reliability.
  - Reduced full batch index (5B+ docs) build time from 2 months to 10 days.
  - Reduced real-time feeding system database queries by 60%.

Flickr / Yahoo! | Senior Operations Engineer | 2005-2010
- Built monitoring and metric collection software to ensure maximum uptime for flickr.com.
- Reduced database logging 1000 qps with spread/apr streaming log consumer of real-time Flickr API usage.
- Automated deployment, taking bare metal to serving applications for thousands of machines.

Various Software Development Positions with Accenture and Comcast | 2001-2005

Education
- Various open source contributions - boto/gearman/ganglia/log4j/rrdtool/..., authored s3nb, ipython-elasticsearch
- BS in MSIS from Penn State, 1998-2001, Graduate Certificate in IT from University of Queensland, 2009